


Native input work across cache states — d1-small-reuse24



Each table contributes 1 first-use and 7 subsequent queries; tables repeat across balanced pairs. Later first-use requests can reuse a common instruction prefix. Panels use different y scales. Token work is not latency or an API-price estimate; quality gates still apply.